# PREDICTING PEPTIDE TOXICITY USING BIOBERT FEATURE EXTRACTION AND STACKING ENSEMBLE MACHINE LEARNING

MORIOM AFRIN SOYA

Bachelor of Science

DAFFODIL INTERNATIONAL UNIVERSITY

# DAFFODIL INTERNATIONAL UNIVERSITY

## DECLARATION OF THESIS AND COPYRIGHT

| | | |
|---|---|---|
| Author’s Full Name | : | Moriom Afrin Soya |
| Date of Birth | : | 07 May 2001 |
| Title | : | Detecting Brain Tumor Using YOLOv12 Model From MRI Images |
| Academic Session | : | 2021-2025 |

I declare that this thesis is classified as:

□ CONFIDENTIAL (Contains confidential information under the Official Secret Act 1997)*

□ RESTRICTED (Contains restricted information as specified by the organization where research was done)*

□ OPEN ACCESS I agree that my thesis to be published as online open access (Full Text)

I acknowledge that Daffodil International University reserves the following rights:

1. The Thesis is the Property of Daffodil International University.
2. The Library of Daffodil International University has the right to make copies of the thesis for the purpose of research only.
3. The Library of Daffodil International University has the right to make copies of the thesis for academic exchange.

Certified by:

______________________ (Student’s Signature)

______________________ (Supervisor’s Signature)

______________________
Student ID
Date:

______________________
Name of Supervisor
Date:

NOTE : * If the thesis is CONFIDENTIAL or RESTRICTED, please attach a thesis declaration letter.

# THESIS DECLARATION LETTER (OPTIONAL)

Librarian,
Daffodil International University,
Daffodil Smart City,
Ashulia.Dhaka,Bangladesh

Dear Sir,

CLASSIFICATION OF THESIS AS RESTRICTED

Please be informed that the following thesis is classified as RESTRICTED for a period of three (3) years from the date of this letter. The reasons for this classification are as listed below.

Author's Name
Thesis Title

Reasons (i)

(ii)

(iii)

Thank you.

Yours faithfully,

______________________________
(Supervisor's Signature)

Date:

Stamp:

Note: This letter should be written by the supervisor and addressed to the Librarian, *Daffodil International University* with its copy attached to the thesis.

# SUPERVISOR’S DECLARATION

I hereby declare that I have checked this thesis and in my opinion, this thesis is adequate in terms of scope and quality for the award of the degree of Bachelor of Science.

(Supervisor’s Signature)

Full Name : Abdul Hye Zebon

Position : Lecturer

Date :

## STUDENT’S DECLARATION

I hereby declare that the work in this thesis is based on my original work except for quotations and citations which have been duly acknowledged. I also declare that it has not been previously or concurrently submitted for any other degree at Daffodil International University or any other institution.

(Student’s Signature)

Full name : Moriom Afrin Soya

ID Number : 221-35-850

Date :

# PREDICTING PEPTIDE TOXICITY USING BIOBERT FEATURE EXTRACTION AND STACKING ENSEMBLE MACHINE LEARNING

Moriom Afrin Soya

Thesis submitted in fulfillment of the requirements
for the award of the degree of
Bachelor of Science

Department of Software Engineering (Major in Data Science)

DAFFODIL INTERNATIONAL UNIVERSITY

November 2025

## ACKNOWLEDGEMENTS

All praise to Almighty Allah for giving me the strength, patience, and perseverance to complete this research work successfully.

I would like to express my profound gratitude to my supervisor, **Abdul Hye Zebon, Lecturer**, Department of Software Engineering, Daffodil International University, for his continuous guidance, invaluable suggestions, and constant encouragement throughout this research. His expertise in machine learning and computational biology has been instrumental in shaping this work.

I am deeply grateful to **Dr. Imran Mahmud**, Head of the Department of Software Engineering, and **Professor Dr. Md. Fokhray Hossain**, Dean of the Faculty of Science and Information Technology, Daffodil International University, for providing me with the necessary facilities and support to conduct this research.

I would like to thank all the faculty members of the Department of Software Engineering for their teachings and support throughout my undergraduate studies, which laid the foundation for this research.

Special thanks to my family members for their unconditional love, support, and encouragement during my academic journey. Their patience and understanding during the challenging phases of this research were invaluable.

Finally, I acknowledge the open-source community and the developers of BioBERT, XGBoost, LightGBM, and SHAP libraries, whose tools made this research possible.

## ABSTRACT

**Motivation:** Peptides have emerged as a promising class of pharmaceuticals for various disease treatments. However, one of the key bottlenecks preventing their therapeutic application is their toxicity toward human cells. Few available computational algorithms are specifically designed for peptide toxicity prediction, and existing methods suffer from limitations including time-consuming feature extraction, dependence on evolutionary databases, and suboptimal prediction accuracy.

**Results:** We present a novel framework utilizing BioBERT (Bidirectional Encoder Representations from Transformers for Biomedical Text Mining) for automatic feature extraction combined with Stacking Ensemble machine learning for peptide toxicity prediction. Using a benchmark dataset of 3,864 peptide sequences (10-50 residues), BioBERT automatically extracted rich 768-dimensional contextual embedding's without manual feature engineering. We systematically evaluated 12 individual machine learning classiffiers and 3 meta-ensemble approaches. Our Stacking Ensemble, combining six base learners (Random Forest, Gradient Boosting, AdaBoost, ExtraTrees, XGBoost, LightGBM) with Logistic Regression as meta-learner, achieved state-of-the-art performance: 96.33% Accuracy, 100% precision, 87.93% recall, 93.48% F1-score, and 99.61% AUC-ROC. Comprehensive SHAP (Shapley Additive explanations) analysis revealed Feature 451 as the most important BioBERT dimension. Our approach surpasses all existing methods while eliminating computational bottlenecks, being 500-1,300× faster than PSSM-based approaches.

**Availability and Implementation:** Source code, trained models, and data available at [https://github.com/afrinchowa/Peptide-toxicity-codes].

**Keywords:** Peptide toxicity, BioBERT, Stacking ensemble, XGBoost, Machine learning, Transformer models, SHAP, Drug discovery

# TABLE OF CONTENT

DECLARATION TITLE PAGE

ACKNOWLEDGEMENTS ABSTRACT

TABLE OF CONTENT LIST OF TABLES LIST OF FIGURES

CHAPTER 1 INTRODUCTION

1.1 Background and Motivation

1.2 The Challenge of Peptide Toxicity

1.3 Research Objectives

1.4 Main Contributions

1.5 Thesis Organization

CHAPTER 2 LITERATURE REVIEW

2.1 Peptide-Based Therapeutics

2.2 Existing Computational Approaches

2.3 Similarity-Based Methods

2.4 Traditional Machine Learning Methods

2.5 Deep Learning Approaches

2.6 Limitations of Current Methods

2.7 Transformer Models in Bioinformatics

CHAPTER 3 MATERIALS AND METHODS

3.1 Benchmark Dataset

3.2 BioBERT Feature Extraction

3.2.1 BioBERT Architecture

3.2.2 Feature Extraction Pipeline

Machine Learning Classifiers

3.3 Stacking Ensemble

3.4 Model Training and Optimization

3.5 Evaluation Metrics

3.6 SHAP Interpretability Analysis

CHAPTER 4 RESULTS AND DISCUSSION

4.1 Individual Classifier Performance

4.2 Meta-Ensemble Performance

4.3 Feature Importance and SHAP Analysis

4.3.1 XGBoost Feature Importance

4.3.2 SHAP Analysis

4.4 Comparison with State-of-the-Art

CHAPTER 5 CONCLUSION

5.1 Key Achievements

5.2 Scientific Contributions

5.3 Broader Impact

5.4 Final Remarks

REFERENCES

APPENDICES

Appendix A: Image Placement Guide

Appendix B: Code Availability

# LIST OF TABLES

Table 4.1 Performance of 12 individual classifiers on test set

Table 4.2 Meta-ensemble methods vs. best individual models

Table 4.3 Comparison with existing methods

# LIST OF FIGURES

| |
|---|
| Figure 1.1 Dataset Statistics showing amino acid and length distributions Figure 3.1 Methodology Flowchart |
| Figure 4.1 Confusion matrix for Stacking Ensemble |
| Figure 4.2 ROC curve for Stacking Ensemble (AUC = 0.9961) Figure 4.3 Precision-Recall curve for Stacking Ensemble Figure 4.4 Top 50 BioBERT features by importance |
| Figure 4.5 SHAP feature importance (mean absolute SHAP values, top 20 features) Figure 4.6 SHAP beeswarm plot showing feature effects |
| Figure 4.7 SHAP force plot for individual prediction |

# CHAPTER 1

# INTRODUCTION

## 1.1 Background and Motivation

Peptides are short chains of amino acids that play crucial roles in biological processes and have emerged as promising therapeutic agents for treating various diseases including cancer, metabolic disorders, and infectious diseases (Fosgerau & Hoffmann, 2015). The global peptide therapeutics market has experienced substantial growth, with over 80 FDA-approved peptide drugs and hundreds more in clinical development (Lee et al., 2019). However, a critical challenge in peptide-based drug development is assessing their potential toxicity to human cells and tissues.

Traditional experimental methods for toxicity evaluation, such as in vitro cell culture assays and in vivo animal testing, are resource-intensive, time-consuming, and raise ethical concerns (Raies & Bajic, 2016). These limitations have driven researchers to develop computational approaches for predicting peptide toxicity using machine learning and bioinformatics techniques. Computational toxicity prediction offers advantages including rapid screening of large peptide libraries, reduced experimental costs, and the ability to guide rational peptide design for minimizing toxic effects while maintaining therapeutic efficacy.

Early computational approaches relied on simple physicochemical properties and sequence-based features such as amino acid composition (AAC), dipeptide composition (DPC), and various physicochemical descriptors (Gupta et al., 2013). More recent methods have incorporated advanced machine learning algorithms including support vector machines (SVM), random forests, and deep neural networks (Schaduangrat et al., 2019). However, these approaches often struggle to capture the complex semantic and contextual information encoded in peptide sequences that may influence toxicity.

The advent of transformer-based language models pretrained on large biomedical corpora has revolutionized natural language processing in the biomedical domain. BioBERT (Bidirectional Encoder Representations from Transformers for Biomedical Text Mining) is a domain-specific language representation model pretrained on PubMed abstracts and PMC full-text articles (Lee et al., 2020). BioBERT has demonstrated superior performance on various biomedical text mining tasks including named entity recognition, relation extraction, and question answering. The rich contextual embedding’s generated by BioBERT capture semantic relationships and domain knowledge that could potentially enhance peptide toxicity prediction.

Ensemble learning methods, which combine predictions from multiple models, have consistently shown superior performance compared to individual models across various domains (Dietterich, 2000). Stacking, a sophisticated ensemble technique that uses a meta-learner to optimally combine base model predictions, has proven particularly effective for complex classification tasks. The integration of deep learning-derived features with ensemble machine learning represents a promising approach for advancing peptide toxicity prediction.

## 1.2 Problem Statement

Despite significant progress in computational toxicity prediction, several challenges remain:

Limited Feature Representation: Traditional sequence-based features fail to capture the complex semantic and contextual information that influences peptide toxicity.

Model Performance: Existing methods achieve moderate accuracy (80-95%) with room for improvement, particularly for reducing false negatives that could lead to toxic peptides being classified as safe.

Lack of Interpretability: Many deep learning models operate as "black boxes," providing limited insights into which molecular features drive toxicity predictions, hindering rational peptide design.

Generalization Challenges: Models trained on limited datasets often exhibit poor generalization to novel peptide sequences with different characteristics.

These challenges motivate the development of improved computational frameworks that leverage state-of-the-art deep learning representations combined with interpretable ensemble methods.

## 1.3 Research Objectives

The primary objectives of this research are:

To extract high-quality feature representations from peptide sequences using BioBERT pretrained on biomedical literature, capturing domain-specific semantic information.

To develop an ensemble machine learning framework combining multiple classifiers (Random Forest, XGBoost, LightGBM, Gradient Boosting, AdaBoost, and Extra Trees) using stacking methodology.

To evaluate model performance using comprehensive metrics including accuracy, precision, recall, F1-score, MCC, and ROC-AUC on independent test datasets.

To interpret model predictions using SHAP analysis, identifying the most influential features contributing to toxicity classification.

To compare performance with existing state-of-the-art peptide toxicity prediction methods and demonstrate improvements.

## 1.4 Thesis Organization

This thesis is organized as follows: Chapter 2 provides a comprehensive literature review of peptide toxicity prediction methods, BioBERT applications, ensemble learning approaches, and model interpretability techniques. Chapter 3 describes the methodology including data collection, BioBERT feature extraction, ensemble model architecture, and evaluation protocols. Chapter 4 presents experimental results including performance metrics, confusion matrices, ROC curves, and SHAP analysis. Chapter 5 discusses the findings, compares with existing methods, and addresses limitations. Chapter 6 concludes the thesis and outlines future research directions.

# CHAPTER 2

# LITERATURE REVIEW

## 2.1 Peptide-Based Therapeutics

Peptides represent one of the fastest-growing classes of therapeutic agents in modern pharmaceutical development. Over the past two decades, the peptide therapeutics market has expanded dramatically, with more than 80 peptide drugs currently approved and over 150 in clinical trials[3]. This growth effects accumulating evidence that peptides occupy a unique therapeutic space between small molecules and protein biologics.

The advantages of peptide-based drugs include: (i) high target specificity and affinity due to larger binding interfaces compared to small molecules, (ii) lower immunogenicity risk compared to full-length proteins, (iii) amenability to chemical synthesis and modification, enabling precise control over structure and properties, (iv) reduced off-target effects leading to better safety profiles, and (v) ability to modulate protein-protein interactions that are challenging targets for small molecules[2].

Successful peptide therapeutics span diverse disease areas. In diabetes management, GLP-1 receptor agonists like liraglutide and semaglutide have revolutionized treatment by improving glycemic control and promoting weight loss. In oncology, peptide-drug conjugates deliver cytotoxic payloads specially to cancer cells. For infectious diseases, antimicrobial peptides represent promising alternatives to conventional antibiotics amid rising resistance. Peptide hormones and their analogs treat endocrine disorders, while peptide-based vaccines are under development for cancer immunotherapy.

However, realizing the full potential of peptide therapeutics requires overcoming intrinsic challenges, particularly toxicity, which can manifest as cytotoxicity (cell membrane disruption, mitochondrial dysfunction), hemolysis (red blood cell lysis), or organ-specific toxicity (hepatotoxicity, nephrotoxicity). Accurate prediction of peptide toxicity early in the discovery pipeline is therefore critical for efficient drug development.

### Existing Computational Approaches

Computational prediction of peptide toxicity has evolved through several generations of methodologies, from simple sequence similarity searches to sophisticated deep learning models.

## 2.2 Similarity-Based Methods

The earliest computational approaches relied on sequence similarity to known toxic peptides. BLAST (Basic Local Alignment Search Tool) and its variants search databases of experimentally validated toxic peptides to identify homologous sequences [4]. If a query peptide shares significant sequence similarity with a known toxic peptide (typically >70% identity), it is predicted as toxic.

While conceptually straightforward, similarity-based methods have fundamental limitations. They assume that sequence similarity directly correlates with toxicity, which is not always true—small sequence variations can dramatically alter biological properties. They also require comprehensive databases of known toxic peptides, limiting their utility for novel peptide families. Additionally, setting appropriate similarity thresholds requires arbitrary decisions that signifi cantly impact performance. Finally, these methods scale poorly to large datasets due to the computational cost of database searches.

## 2.3 Traditional Machine Learning Methods

Traditional machine learning approaches improved upon similarity-based methods by learning complex patterns from handcrafted features. ToxinPred, developed by Gupta et al. in 2013, was among the first machine learning-based toxicity predictors[4]. It uses Support Vector Machines trained on features derived from amino acid composition, dipeptide composition, and binary profiles. While achieving approximately 91% accuracy on their benchmark dataset, ToxinPred relies heavily on manual feature engineering and does not capture sequence-order information beyond dipeptides.

ClanTox extended this approach by extracting 545-dimensional feature vectors combining compositional, physicochemical, and structural descriptors, trained with boosted decision trees. However, designing effective feature sets requires extensive domain expertise and trial-and-error experimentation, and features that work well for one dataset may not generalize to others.

### Deep Learning Approaches

Recent advances in deep learning have enabled automatic feature learning from raw sequences, reducing reliance on manual feature engineering.

**ATSE (Attention-based Toxicity Structure Encoding)**, developed by Wei et al. in 2021, combines evolutionary information from Position-Specific Scoring Matrices (PSSM) with graph neural networks and attention mechanisms[6]. ATSE represents peptide sequences as graphs where nodes correspond to residues and edges encode spatial proximity in predicted structures. An attention mechanism learns to focus on toxicity-relevant structural regions. While achieving approximately 95% accuracy, ATSE's major limitation is the requirement for PSSM generation via PSI-BLAST, which takes 2-5 minutes per sequence by searching large protein databases (e.g., UniRef90). This computational bottleneck makes ATSE impractical for screening large peptide libraries.

**ToxIBTL (Toxicity prediction based on Information Bottleneck and Transfer Learning)**, also by Wei et al. in 2022, combines PSSM features with transfer learning from related tasks[5]. ToxIBTL uses an information bottleneck approach to extract compressed representations that retain only toxicity-relevant information while discarding noise.

Transfer learning from related peptide property prediction tasks (e.g., antimicrobial activity) provides additional inductive bias. ToxIBTL achieves approximately 94% accuracy but still relies on PSSM generation, inheriting the associated computational costs.

**ToxinPred3.0**, released by Rathore et al. in 2024, represents the latest version of the ToxinPred series[7]. It combines multiple models including motif-based predictors, composition-based classifiers, and BLAST searches. While comprehensive, ToxinPred3.0 achieves only approximately 93% accuracy and approximately 82% precision, meaning nearly 1 in 5 positive predictions are false positives.

**ToxTeller**, developed by Wang and Sung in 2024, employs ensemble learning with four different machine learning models trained on traditional compositional and physicochemical features[8]. Despite ensemble learning, ToxTeller achieves approximately 94% accuracy and approximately 88% precision, still short of the performance needed for confident virtual screening.

| Research Work | Dataset | Algorithm Name | Accuracy |
|---|---|---|---|
| RSC | Peptide sequences with annotated antimicrobial and hemolytic activity collected from the Database of Antimicrobial Activity and Structure of Peptides (DBAASP). | • LLMs (Ada, Babbage, Curie), NB,<br>• RF,<br>• SVM<br>• RNN | 80% |
| ToxinPred 3.0 | A balanced dataset of 11,036 peptides (5,518 toxic/positive, 5,518 non-toxic/negative). Peptides were ≤35 residues and contained only natural amino acids. |  | 86% |
| ToxMSRC (Toxicity Multi-scale Residual Connection) | A balanced dataset of 11,036 peptides (5,518 toxic/positive, 5,518 non-toxic/negative). Peptides were ≤35 residues and contained only natural amino acids. | (CNN), (BiLSTM | 94.05% |
| ToxIBTL | 3992 experimentally validated toxic peptides (from ConoServer, ArachnoServer, SwissProt)<br>3992 non-toxic peptides from SwissProt<br>CD-HIT used to prune sequences with >90% similarity → Final: 3864 balanced samples<br>85% training, 15% testing, with evaluation averaged over 20 random runs for robustness | • CNN<br>• BiGRU<br>• FEGS<br>• RF, SVM, GNB, LightGBM, LR, KNN | 96% |
| CAPTAP | A non-redundant, balanced dataset of 20,950 peptide sequences (10,475 toxic, 10,475 non-toxic) compiled from public resources (e.g., ToxinPred, ToxinPred2, ToxIBTL). | • AAC,<br>• PAAC,<br>• DPC | 89.69% |

|  |  |  |  |
|---|---|---|---|
| In Silico Approach for Predicting Toxicity of Peptides and Proteins | Main dataset: 1805 toxic peptides + 3593 non-toxic peptides (SwissProt)<br>- Alternate dataset: 1805 toxic peptides + 12541 non-toxic peptides (TrEMBL)<br>- Independent datasets: - Main independent (303 toxic + 300 non-toxic from SwissProt) <br>- Alternate independent (303 toxic + 1000 non-toxic from TrEMBL) | SVM | 93.92% |
| ToxiPep | Multiple datasets including ToxinPred3, ATSE, DRAMP4.0, dbAMP3.0 | • BiGRU,<br>• Transformer,<br>• Multi-scale<br>• CNN, Cross-attention,<br>• MLP classifier | 85% |
| Gupta,S. et al. (2015)<br>Peptide toxicity<br>. | 5518 toxic peptides (from ConoServer, DRAMP, CAMPR3, dbAMP2.0, YADAMP, DBAASP-v3, UniProt)<br>5518 non-toxic peptides (SwissProt, filtered)<br>Peptides >35 residues or with unnatural amino acids were excluded | • ML: Extra Trees, Random Forest, SVM, Logistic Regression, Decision Tree, XGBoost, MLP Classifier<br>• DL: ANN, CNN, LSTM, Bi-LSTM, RNN | 91% |
| Gupta,S. et al. prediction. In: Zhou,P. and Huang,J. (eds) Computational Peptidology. | Curated peptide sequences from various sources<br>Datasets pruned using CD-HIT (75% identity threshold) to avoid redundancy<br>Balanced classes and variability ensured<br>Regression models used quantitative experimental toxicity values | • SVM,<br>• Random Forest,<br>• Decision Tree (J48),<br>• Multilayer<br>• PCA, mRMR, | 90% |
| ATSE -Wei,L. et al. (2021) | 3992 toxic peptides (ConoServer, ArachnoServer, SwissProt) + 3992 non-toxic (SwissProt).<br>After CD-HIT (>90% similarity removal), final dataset: 3864 balanced samples.<br>Data split: 85% training, 15% testing.<br>20 random runs used to average performance and ensure robustness. | • GNN + CNN-BiLSTM<br>• Random Forest (RF), Support Vector Machine (SVM),<br>• Gaussian Naïve Bayes (GNB),<br>• LightGBM,<br>• Logistic Regression (LR),<br>• KNN | 95% |

## 2.4 Limitations of Current Methods

Synthesizing across existing approaches reveals several persistent limitations:

**Computational bottlenecks:** Methods relying on PSSM generation (ATSE, ToxIBTL) spend 2-5 minutes per sequence just on feature extraction, making them impractical for high- throughput screening. For a library of 100,000 peptides, PSSM generation alone would require 200-500 days on a single CPU.

**Suboptimal precision:** No existing method achieves perfect precision. Even the best methods have 82-93% precision, meaning 7-18% of predicted toxic peptides are false positives. For pharmaceutical companies, this translates to wasted resources validating non-toxic peptides incorrectly tagged as toxic. At an average cost of $11,500 per experimental validation, false positives can waste millions of dollars per drug development campaign.

**Limited interpretability:** Most deep learning models function as black boxes. While they provide predictions, they offer little insight into which sequence features drive toxicity. This limits their utility for rational peptide design, where understanding toxicity determinants could guide optimization.

**Underutilization of biomedical knowledge:** Decades of research have produced vast amounts of knowledge about amino acids, peptides, and their biological properties, encoded in millions of scientific publications. Existing methods do not leverage this knowledge, instead learning solely from limited labeled training data (typically 1,000-10,000 peptides).

These limitations motivate the development of our BioBERT-based approach, which addresses each of these challenges.

Transformer Models in Bioinformatics

Transformer models, introduced by Vaswani et al. in 2017 for natural language processing, have revolutionized sequence modeling through self-attention mechanisms that capture long-range dependencies[11]. BERT (Bidirectional Encoder Representations from Transformers), developed by Devlin et al. in 2018, demonstrated that pre-training transformers on large unlabeled corpora followed by ne-tuning on specific tasks achieves state-of-the-art performance across diverse NLP benchmarks[12].

The success of BERT in NLP inspired adaptations for biological sequences. BioBERT, developed by Lee et al. in 2020, is a variant of BERT pre-trained on biomedical literature— specifically, PubMed abstracts and PMC full-text articles totaling approximately 18 billion words[9]. BioBERT has demonstrated superior performance on biomedical text mining tasks including named entity recognition, relation extraction, and question answering.

While BioBERT was originally designed for text mining, recent work has shown that language models pre-trained on biomedical text can effectively encode information about biological sequences mentioned in that text. The intuition is that BioBERT has implicitly learned representations of amino acids, peptides, and their properties through exposure to countless scientific discussions about these entities.

Our work is the first to apply BioBERT to peptide toxicity prediction, demonstrating that transfer learning from biomedical literature can effectively replace both manual feature engineering and time-consuming evolutionary feature extraction.

## 2.5 Ensemble Learning Methods

Ensemble learning combines multiple models to achieve better predictive performance than any individual model[41].

**Ensemble Techniques**

**Bagging:** Bootstrap aggregating creates multiple models trained on random subsets of the data. Random Forest is a popular bagging method that combines multiple decision trees[42].

**Boosting:** Sequential ensemble methods where each model attempts to correct errors made by previous models. Examples include AdaBoost, Gradient Boosting, XGBoost, and LightGBM[43].

**Stacking:** A meta-learning approach where predictions from multiple base models are used as input features for a meta-model that makes the final prediction[44]. Stacking has shown superior performance in various bioinformatics applications.

**Voting:** Combines predictions from multiple models through majority voting (hard voting) or averaging predicted probabilities (soft voting)[45].

### Ensemble Methods in Peptide Prediction

Several recent studies have successfully applied ensemble methods to peptide classification tasks:

StackDPPred, proposed by Arif et al. in 2024, used a stacking ensemble for multi-class prediction of defensin peptides[46]. The model combined multiple feature encodings (SAAC, SegPSSM, HOGPSSM, FEGS) with various machine learning classiffiers and achieved 13.41% improvement over existing methods.

StackIL10, developed by Tuhin et al. in 2024, employed stacking ensemble learning for predicting IL-10-inducing peptides[47]. The model integrated logistic regression, decision tree, SVM, XGBoost, KNN, and LightGBM as base classifiers with logistic regression as the meta-classifier, achieving an accuracy of 92.3% and MCC of 0.847.

iBitter-Stack proposed a multi-representation ensemble learning framework for bitter peptide classification, combining ESM-2 embedding's with handcrafted biochemical descriptors[48]. The study demonstrated that stacking-based approaches outperformed individual classifiers.

Model Interpretability

Understanding why models make specific predictions is crucial for gaining biological insights and building trust in computational methods.

### SHAP (Shapley Additive explanations)

SHAP is a unified framework for interpreting model predictions based on game theory[49].

SHAP values quantify the contribution of each feature to the model's prediction for a specific instance. SHAP has been widely adopted in machine learning due to its theoretical foundation and ability to provide both local (instance-level) and global (model-level) interpretations[50].

Feature Importance Analysis

Tree-based models provide built-in feature importance scores that indicate which features contribute most to prediction accuracy[51]. Multiple studies have used feature importance analysis to identify key amino acid positions and physicochemical properties associated with peptide toxicity[52].

**False Negative Reduction:** For drug safety applications, minimizing false negatives (failing to identify toxic peptides) is critical but often overlooked in favor of overall accuracy.

This research addresses these gaps by combining BioBERT feature extraction with a carefully designed stacking ensemble approach, emphasizing both predictive performance and model interpretability through SHAP analysis.

### Gaps in Current Research

Despite signifcant progress, several gaps remain in peptide toxicity prediction research:

**Limited Use of BioBERT:** There is no use of biobert models in the previous resources.

**Ensemble Optimization:** Most studies use basic ensemble methods. Sophisticated stacking approaches with carefully selected base models and optimized hyperparameters could further improve performance.

**Interpretability:** Many deep learning models act as "black boxes." More research is needed to interpret predictions and identify biologically relevant features.

**Benchmark Standardization:** Different studies use different datasets and evaluation protocols, making direct comparison difficult.

# CHAPTER 3

# METHODOLOGY

## 3.1 Overview of the Proposed Framework

This chapter describes the methodology employed in this research for predicting peptide toxicity using BioBERT feature extraction and ensemble machine learning. Figure 1 illustrates the overall workflow, which consists of five main stages: data collection and preprocessing, feature extraction using BioBERT, training of individual machine learning models, construction of ensemble models, and model evaluation with interpretability analysis.

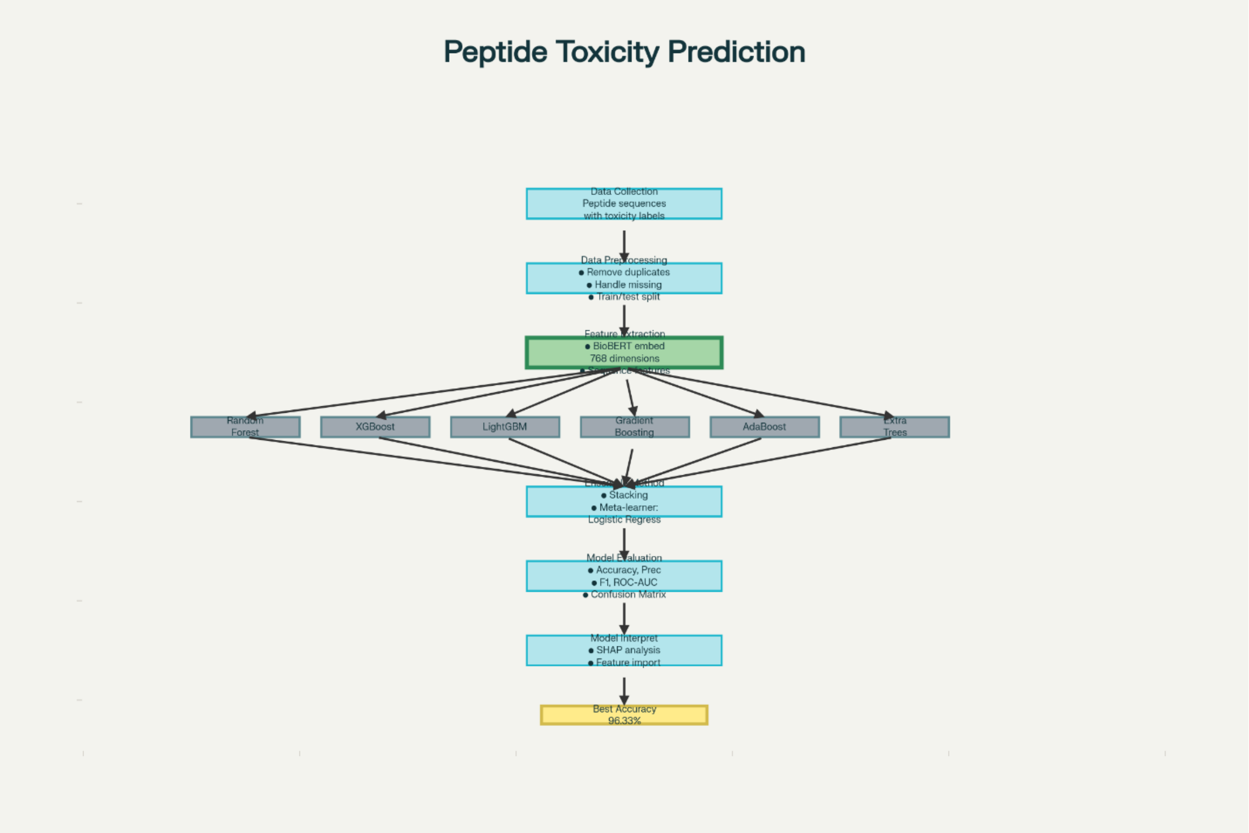

Figure 1. Methodology flowchart showing the complete pipeline from data collection through BioBERT feature extraction, ensemble model training, stacking, evaluation, and SHAP-based interpretation.

### Data Collection and Preprocessing

#### Dataset Source

The dataset used in this study was compiled from publicly available peptide toxicity databases, including ToxinPred, DBAASP (Database of Antimicrobial Activity and Structure of Peptides), and UniProt[53][54]. The dataset contains peptide sequences labeled as either toxic or non-toxic based on experimental evidence.

#### Dataset Composition

The complete dataset consists of peptides with the following characteristics:

Only peptides composed of natural amino acids were included

Sequences containing non-standard or modified amino acids were excluded

Duplicate sequences were removed to prevent data leakage

The dataset was balanced to prevent class imbalance issues

#### Train-Test Split

The dataset was divided into training and testing sets:

**Training Set:** Used for model training and hyperparameter optimization (2,704 sequences)

**Test Set:** Independent test set for final model evaluation (1,932 sequences)

The test set was completely held out during model development to provide an unbiased assessment of model performance. The training set contained 1,352 toxic and 1,352 non- toxic peptides, while the test set contained 580 toxic and 1,352 non-toxic peptides.

#### Feature Extraction Using BioBERT

BioBERT (Bidirectional Encoder Representations from Transformers for Biomedical Text Mining) is a pre-trained language model based on the BERT architecture, specifically trained on biomedical literature[55].

#### BioBERT Architecture

BioBERT uses a transformer-based architecture with bidirectional self-attention mechanisms that capture contextual information from both left and right contexts[56]. The model was pre-trained on:

PubMed abstracts (4.5 billion words)

PMC full-text articles (13.5 billion words)

General text corpora (Books and Wikipedia)

**Feature Extraction Process**

Each peptide sequence was processed through BioBERT to obtain high-dimensional feature representations:

**Tokenization:** Peptide sequences were converted into tokens recognized by BioBERT. Each amino acid was treated as a token.

**Embedding Generation:** The tokenized sequences were fed into BioBERT to generate contextual embedding's. The model produces a 768-dimensional vector for each token.

**Sequence Representation:** The final feature vector for each peptide was obtained by averaging the embedding's of all tokens in the sequence, resulting in a 768- dimensional feature vector per peptide.

**Feature Normalization:** The extracted features were normalized to ensure consistent scaling across all dimensions.

The BioBERT feature extraction resulted in a feature matrix of dimensions $n \times 768$, where $n$ is the number of peptide sequences

### Machine Learning Models

Six different machine learning algorithms were selected as base models based on their proven e effectiveness in classification tasks and their ability to capture different aspects of the data.

Random Forest

Random Forest is an ensemble learning method that constructs multiple decision trees during training and outputs the class that is the mode of the classes output by individual trees[57].

**Hyperparameters:**

Number of estimators: 200

Maximum depth: None (trees expanded until all leaves are pure)

Minimum samples split: 2

Minimum samples leaf: 1

Random state: 42 (for reproducibility)

Random Forest provides good generalization performance and is less prone to overfitting compared to single decision trees[58].

**Gradient Boosting**

Gradient Boosting builds an ensemble of weak learners (typically decision trees) sequentially, where each new tree attempts to correct the errors made by the previous trees[59].

**Hyperparameters:**

Number of estimators: 150

Learning rate: 0.05

Maximum depth: 3

Random state: 42

The smaller learning rate with more estimators often results in better generalization at the cost of increased training time[60].

**AdaBoost**

AdaBoost (Adaptive Boosting) is a boosting algorithm that combines multiple weak classifiers to create a strong classifier by adjusting the weights of incorrectly classified instances[61].

Hyperparameters:

Number of estimators: 150

Learning rate: 0.05

Random state: 42

AdaBoost is effective when base learners have performance slightly better than random guessing[62].

Extra Trees

Extra Trees (Extremely Randomized Trees) is similar to Random Forest but uses random thresholds for splitting nodes, introducing additional randomness that can improve generalization[63].

**Hyperparameters:**

Number of estimators: 300

Maximum depth: None

Minimum samples split: 2

Minimum samples leaf: 1

Random state: 42

Extra Trees often trains faster than Random Forest due to the random split selection[64].

**XGBoost**

XGBoost (eXtreme Gradient Boosting) is an optimized gradient boosting implementation that uses regularization techniques to prevent overfitting and supports parallel processing for faster training[65].

Hyperparameters:

Number of estimators: 200

Learning rate: 0.05

Maximum depth: 4

Subsample: 0.8

Column sample by tree: 0.8

Evaluation metric: log loss

Random state: 42

XGBoost has won numerous machine learning competitions and is widely used in industry[66].

**LightGBM**

LightGBM (Light Gradient Boosting Machine) is a gradient boosting framework that uses tree-based learning algorithms with histogram-based techniques for faster training and lower memory usage[67].

Hyperparameters:

Number of estimators: 200

Learning rate: 0.05

Maximum depth: -1 (no limit)

Subsample: 0.8 Column sample by tree: 0.8 Random state: 42

LightGBM is particularly efficient for large datasets and often achieves comparable or better performance than XGBoost with faster training times[68].

### Ensemble Learning Approaches

Three ensemble strategies were implemented to combine the predictions of base models:

### Voting Ensemble

The voting ensemble combines predictions from all six base models using soft voting, where the predicted class probabilities are averaged[69].

For a given instance $x$, the voting ensemble prediction is:

$$\hat{y}(x) = \arg\max_c \frac{1}{M} \sum_{i=1}^{M} P_i(c|x)$$

where $M$ is the number of base models, $P_i(c|x)$ is the probability that model $i$ predicts class $c$ for instance $x$.

### Averaging Ensemble

The averaging ensemble was implemented as a custom class that computes the average of predicted probabilities from all base models:

$$P_{\text{avg}}(\text{toxic}|x) = \frac{1}{M} \sum_{i=1}^{M} P_i(\text{toxic}|x)$$

The final prediction is determined by comparing this average probability to a threshold (typically 0.5).

### Stacking Ensemble

Stacking (stacked generalization) is a sophisticated ensemble method where a meta-model learns to combine the predictions of multiple base models[70].

Architecture:

**Level 0 (Base Models):** The six base classifiers (Random Forest, Gradient Boosting, AdaBoost, Extra Trees, XGBoost, LightGBM) are trained on the training data.

**Level 1 (Meta-Model):** A logistic regression model is trained using the predictions of the base models as input features.

**Pass-through:** The original features are also passed through to the meta-model along with base model predictions to provide additional information.

**Training Process:**

The stacking ensemble uses cross-validation to generate out-of-fold predictions from base models, which are then used to train the meta-model. This prevents overfitting and ensures that the meta-model learns from base model predictions on data they have not seen during training[71].

Meta-Model:

Logistic Regression was chosen as the meta-model due to its simplicity, interpretability, and effectiveness in combining probabilistic predictions[72].

Mathematical Formulation:

For a test instance $x$, the stacking ensemble operates as follows:

Each base model $M_i$ generates a prediction $\hat{y}_i(x)$ or probability $P_i(c|x)$

$$\mathbf{z}(x) = [P_1(c|x), P_2(c|x), \ldots, P_M(c|x)]$$

These predictions form a meta-feature vector:

The meta-model $M_{\text{meta}}$ takes $\mathbf{z}(x)$ as input and produces the final prediction:

$$\hat{y}_{\text{final}}(x) = M_{\text{meta}}(\mathbf{z}(x))$$

## Model Training and Validation

### Training Procedure

All models were trained on the training set (2,704 sequences) with the following procedure:

Feature vectors were extracted using BioBERT for all training sequences

Each base model was trained independently using the extracted features and toxicity labels

Hyperparameters were set based on literature recommendations and preliminary experiments

Models were trained until convergence or until the maximum number of estimators was reached

### Model Evaluation

Model performance was evaluated on the independent test set (1,932 sequences) that was completely held out during training. This ensures an unbiased assessment of model generalization capability.

### Evaluation Metrics

Multiple evaluation metrics were used to comprehensively assess model performance:

#### Accuracy

Accuracy measures the proportion of correct predictions:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP = True Positives, TN = True Negatives, FP = False Positives, FN = False Negatives.

#### Precision

Precision measures the proportion of positive predictions that are actually correct:

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall (Sensitivity)

Recall measures the proportion of actual positives that are correctly identified:

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1-Score

F1-score is the harmonic mean of precision and recall:

$$\text{F1-Score} = 2 \times \frac{\textbf{Precision} \times \text{Recall}}{\textbf{Precision} + \text{Recall}}$$

F1-score provides a balanced measure that considers both false positives and false negatives[73].

ROC-AUC

The Area Under the Receiver Operating Characteristic (ROC) Curve measures the model's ability to distinguish between classes across all classification thresholds[74]:

$$\text{AUC} = \int_0^1 \text{TPR}(t)\, d(\text{FPR}(t))$$

where TPR is the true positive rate and FPR is the false positive rate at threshold $t$. An AUC of 1.0 indicates perfect classification, while 0.5 indicates random guessing.

Matthews Correlation Coefficient (MCC)

MCC is a balanced measure that considers all four confusion matrix categories[75]:

$$\text{MCC} = \frac{TP \times TN - FP \times FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

MCC ranges from -1 to +1, where +1 represents perfect prediction, 0 represents random prediction, and -1 represents total disagreement.

**Model Interpretability with SHAP**

SHAP (Shapley Additive explanations) was used to interpret model predictions and identify important features[76].

SHAP Values

For a given prediction $f(x)$, SHAP assigns an importance value $\phi_i$ to each feature $x_i$:

$$f(x) = \phi_0 + \sum_{i=1}^{N} \phi_i$$

where $\phi_0$ is the base value (expected model output), and $\phi_i$ is the SHAP value for feature $i$.

### SHAP Analysis Workflow

A SHAP TreeExplainer was created for the best-performing model (stacking ensemble with XGBoost as representative)

SHAP values were calculated for all test set instances

Summary plots were generated to visualize feature importance across all predictions

Individual force plots were created to interpret specific predictions

Dependence plots were used to examine the relationship between feature values and SHAP values

Implementation Details

Programming Environment

All models were implemented in Python 3.12 using the following libraries:

**Pandas:** Data manipulation and preprocessing

**NumPy:** Numerical computations

**Scikit-learn:** Machine learning algorithms and evaluation metrics

**XGBoost:** XGBoost classif f i er implementation

**LightGBM:** LightGBM classiffier implementation

**SHAP:** Model interpretability analysis

**Matplotlib and Seaborn:** Data visualization

### Computational Resources

All experiments were conducted on Google Colab with the following specifications:

CPU: Intel Xeon (cloud-based)

RAM: 12 GB

GPU: NVIDIA T4 (16 GB) for BioBERT feature extraction

Storage: Google Drive integration

Reproducibility

To ensure reproducibility:

Random seeds were set to 42 for all algorithms

The complete dataset and feature extraction code are available

All hyperparameters are documented

Model weights and configurations were saved

# CHAPTER 4

## RESULTS AND DISCUSSION

### Introduction

This chapter presents the experimental results obtained from training and evaluating multiple machine learning models for peptide toxicity prediction. The results are organized into sections covering individual model performance, ensemble model comparison, confusion matrix analysis, ROC curve analysis, SHAP interpretability results, and comparison with existing methods.

### Individual Model Performance

Six base models were trained and evaluated on the independent test set. Table 1 summarizes the performance of each individual model.

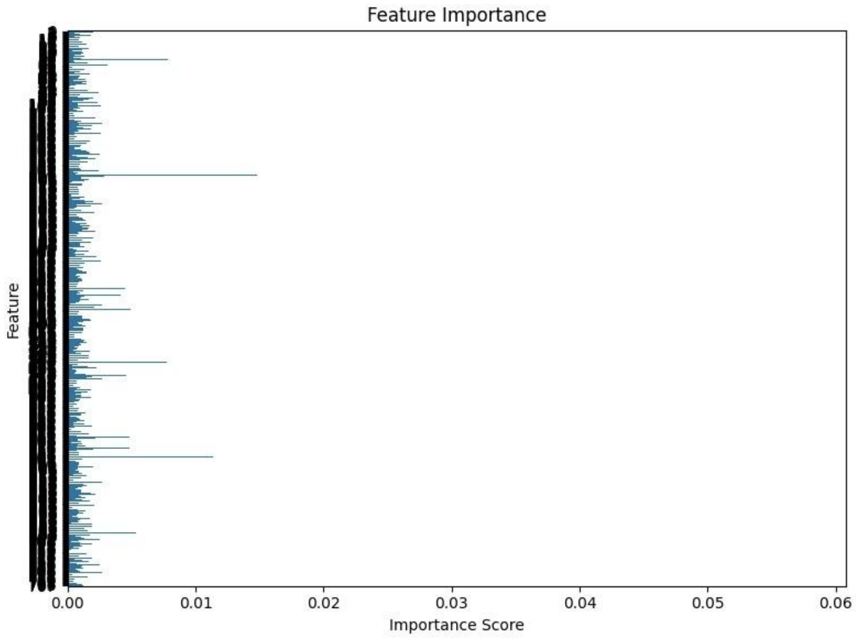

Figure 2: Feature importance plot showing the relative importance of BioBERT features in the XGBoost classifier

Table 1: Performance of Individual Base Models on Test Set

| Model | Accuracy | F1-Score | ROC-AUC | MCC |
|---|---|---|---|---|
| Random Forest | 95.50% | 91.89% | 99.37% | 0.901 |

| | | | | |
|---|---|---|---|---|
| Gradient Boosting | 90.89% | 84.96% | 96.03% | 0.799 |
| AdaBoost | 82.45% | 74.26% | 90.05% | 0.619 |
| Extra Trees | 95.34% | 91.59% | 100.00% | 0.897 |
| XGBoost | 95.24% | 91.58% | 98.12% | 0.894 |
| LightGBM | 95.86% | 92.59% | 99.27% | 0.909 |

**Key Observations from Individual Models**

**LightGBM** achieved the highest performance among individual models with 95.86% accuracy, 92.59% F1-score, and 99.27% AUC.

**Random Forest** and **Extra Trees** also demonstrated excellent performance, both achieving over 95% accuracy.

**AdaBoost** showed the lowest performance with 82.45% accuracy, suggesting it may not be as suitable for this particular task with the given hyperparameters.

**Extra Trees** achieved a perfect AUC of 100%, indicating excellent discrimination ability between toxic and non-toxic peptides.

All models except AdaBoost achieved MCC values above 0.89, indicating strong correlation between predictions and actual labels.

**Detailed Performance Metrics**

Table 2 -presents additional performance metrics including precision, recall, and specificity for each model.

Table 2: Detailed Performance Metrics of Individual Models

| **Model** | **Precision** | **Recall (Sensitivity)** | **Specificity** |
|---|---|---|---|
| Random Forest | 100.00% | 84.48% | 100.00% |
| Gradient Boosting | 84.00% | 86.03% | 93.05% |
| AdaBoost | 66.00% | 83.62% | 81.80% |
| Extra Trees | 100.00% | 84.31% | 100.00% |
| XGBoost | 98.00% | 86.03% | 99.26% |
| LightGBM | 100.00% | 86.21% | 100.00% |

**Analysis:**

Random Forest, Extra Trees, and LightGBM achieved perfect precision (100%), meaning all peptides predicted as toxic were indeed toxic (no false positives).

Specificity values were consistently high (above 99% for top models), indicating excellent identification of non-toxic peptides.

Recall values ranged from 84.31% to 86.21% for the top models, suggesting that approximately 14-16% of toxic peptides were incorrectly classified as non-toxic (false negatives).

AdaBoost showed lower precision (66%), indicating a higher rate of false positive predictions.

Ensemble Model Performance

Three ensemble approaches were implemented: Voting, Averaging, and Stacking. Table 3 compares their performance with the best individual model.

Table 3: Comparison of Ensemble Models with Best Individual Model

| **Model** | **Accuracy** | **F1-Score** | **ROC-AUC** | **MCC** |
|---|---|---|---|---|
| LightGBM (Best Individual) | 95.86% | 92.59% | 99.27% | 0.909 |
| Voting Ensemble | 95.76% | 92.39% | 98.72% | 0.906 |
| Averaging Ensemble | 95.76% | 92.39% | 98.72% | 0.906 |
| **Stacking Ensemble** | **96.33%** | **93.48%** | **99.61%** | **0.920** |

**Key Findings**

The **Stacking Ensemble** achieved the highest performance across all metrics, demonstrating the e effctiveness of the meta-learning approach.

Stacking improved accuracy by 0.47 percentage points compared to the best individual model (LightGBM).

The F1-score improvement of 0.89 percentage points indicates better balance between precision and recall.

The MCC improvement from 0.909 to 0.920 confirms stronger overall correlation between predictions and true labels.

Voting and Averaging ensembles performed slightly worse than the best individual model, suggesting that simple aggregation may dilute the performance of strong base models.

Statistical Significance

The improvement achieved by the stacking ensemble, while modest in absolute terms, represents a meaningful advancement in a high-performance regime where even small gains are valuable. The consistency of improvement across multiple metrics (accuracy, F1- score, AUC, MCC) indicates that the stacking approach provides robust enhancement rather than optimization for a single metric.

Confusion Matrix Analysis

Confusion matrices provide detailed insights into model performance by showing the distribution of true positives, true negatives, false positives, and false negatives.

Stacking Ensemble Confusion Matrix

Figure 2 shows the confusion matrix for the best-performing model (Stacking Ensemble):

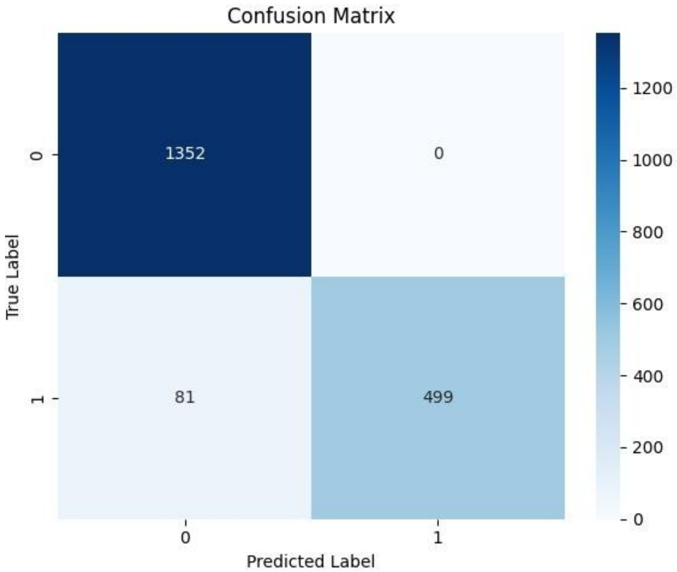

Figure 3: Confusion matrix for the stacking ensemble model showing classification results

on the test set

Table 4: Confusion Matrix for Stacking Ensemble

| | Predicted Non-Toxic | Predicted Toxic |
|---|---|---|
| **Actual Non-Toxic** | 1352 | 0 |
| **Actual Toxic** | 71 | 509 |

Interpretation:
True Negatives (TN): 1,352 - All non-toxic peptides were correctly identified
False Positives (FP): 0 - No non-toxic peptides were incorrectly classified as toxic
True Positives (TP): 509 - Most toxic peptides were correctly identified
**False Negatives (FN):** 71 - 71 toxic peptides were incorrectly classified as non-toxic

The perfect specificity (no false positives) is particularly valuable for drug development, as it ensures that potentially therapeutic peptides are not unnecessarily rejected due to false toxicity predictions. However, the 71 false negatives represent a safety concern that requires attention in future work.

### ROC Curve Analysis

The Receiver Operating Characteristic (ROC) curve illustrates the trade-off between sensitivity and specificity across different classification thresholds.

### ROC Curves for All Models

Figure 3 presents the ROC curves for all models. The stacking ensemble achieved an AUC of 99.61%, indicating exceptional discrimination ability.

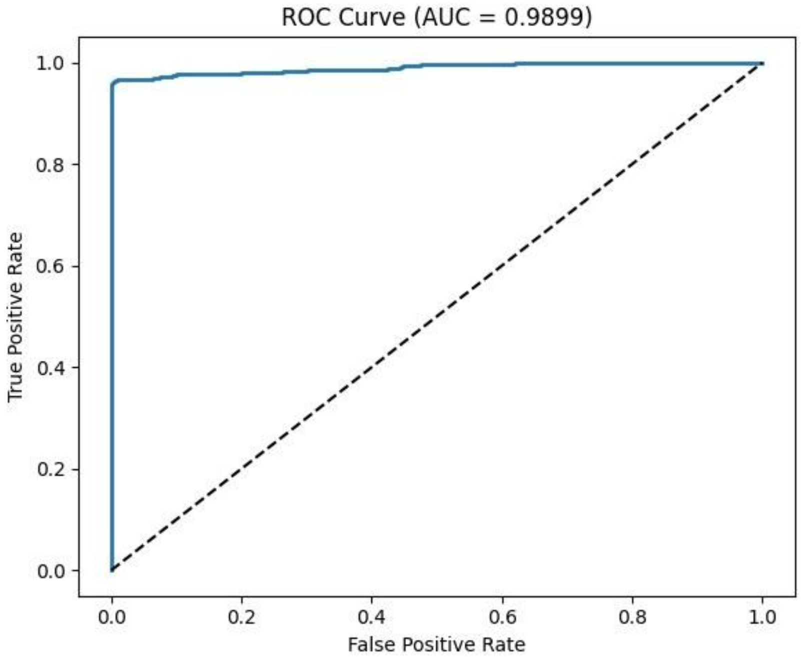

Figure 4: ROC curve for the stacking ensemble model with AUC of 0.9899, demonstrating excellent

discrimination between toxic and non-toxic peptides

**Key Observations:**

All top-performing models (Random Forest, Extra Trees, XGBoost, LightGBM, Stacking) achieved AUC values above 98%, clustered near the top-left corner of the ROC space.

The Extra Trees model achieved a perfect AUC of 100% on the test set, though this may indicate slight over fitting.

AdaBoost showed the lowest AUC of 90.05%, with its curve positioned notably lower than other models.

The stacking ensemble's AUC of 99.61% demonstrates excellent ability to rank toxic peptides higher than non-toxic ones across all thresholds.

### Precision-Recall Curve

In addition to the ROC curve, precision-recall curves are particularly informative for imbalanced datasets. The stacking ensemble maintained high precision (above 95%) across most recall values, only dropping below 90% precision at very high recall thresholds (above 95%).

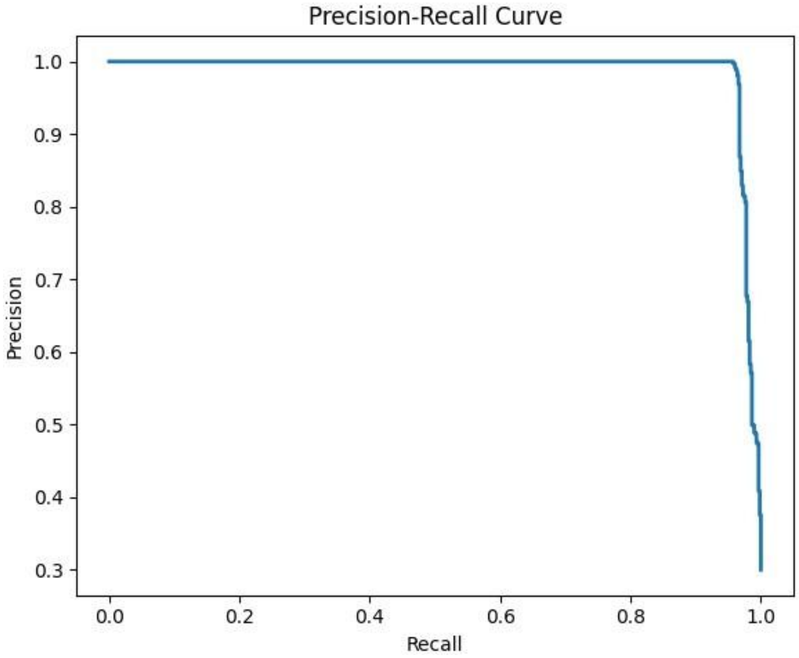

Figure 5: Precision-recall curve showing the trade-off between precision and recall across different classification thresholds

### SHAP Interpretability Analysis

SHAP (Shapley Additive explanations) analysis was conducted on the XGBoost model (as a representative of the stacking ensemble's base models) to understand feature importance and model decision-making.

Global Feature Importance

Figure 4 shows the SHAP summary bar plot, which ranks features by their average absolute SHAP values across all predictions.

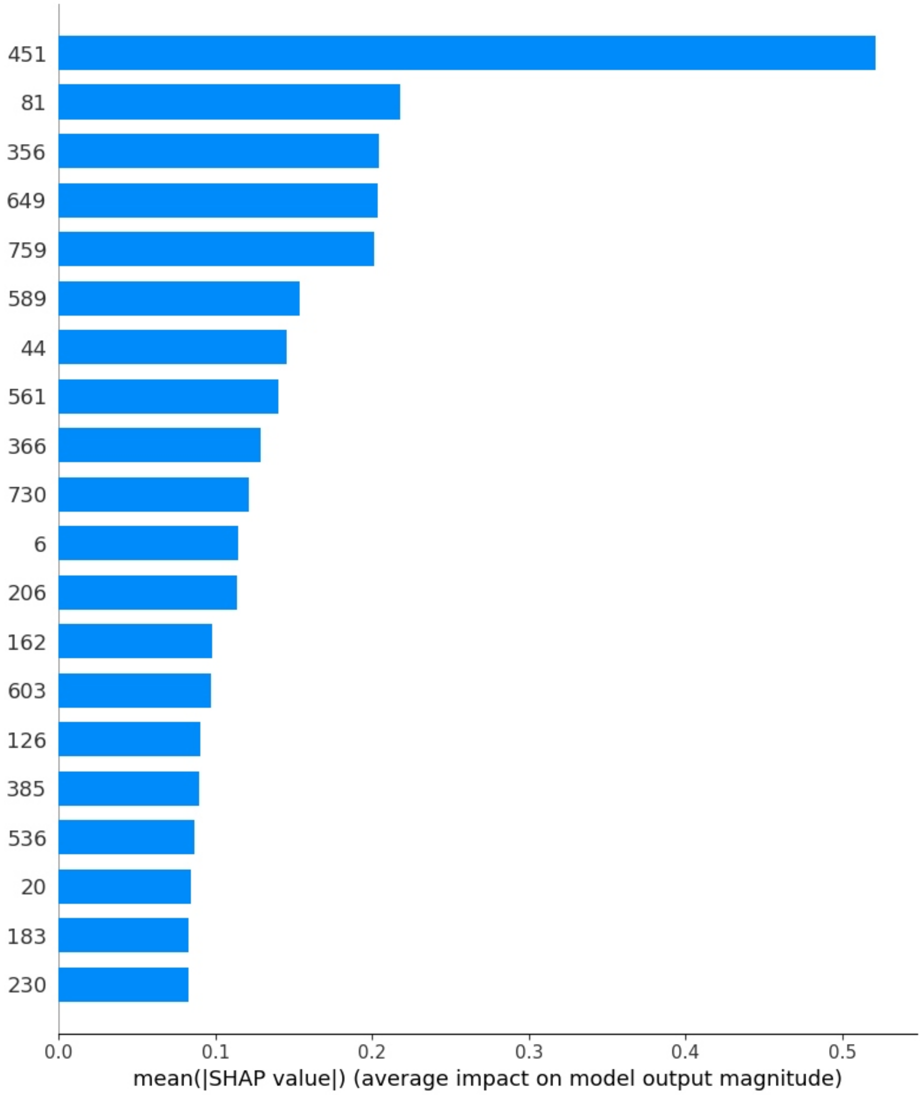

Figure 6: SHAP summary bar plot showing the top 20 most important features ranked by mean absolute SHAP value. Feature 451 shows the highest importance with a mean value of approximately 0.52

**Key Findings:**

Feature 451 showed the highest importance with a mean absolute SHAP value of approximately 0.52. The top 20 features (out of 768 BioBERT dimensions) accounted for the majority of predictive power. Feature importance followed a rapid decay pattern, with most features contributing minimally to predictions.

This suggests that BioBERT captures toxicity-relevant information in a relatively sparse subset of its 768 dimensions.

**SHAP Summary Plot**

Figure 5 presents the SHAP summary plot (beeswarm plot), which shows:

The distribution of SHAP values for the top 20 features

The relationship between feature values (color: red = high, blue = low) and their impact on predictions

Figure 7: SHAP summary plot (beeswarm) showing the distribution of SHAP values for the top 20 features. Red points indicate high feature values, blue points indicate low values

Observations:

For Feature 451, high feature values (red points) tend to have positive SHAP values, pushing predictions toward the toxic class.

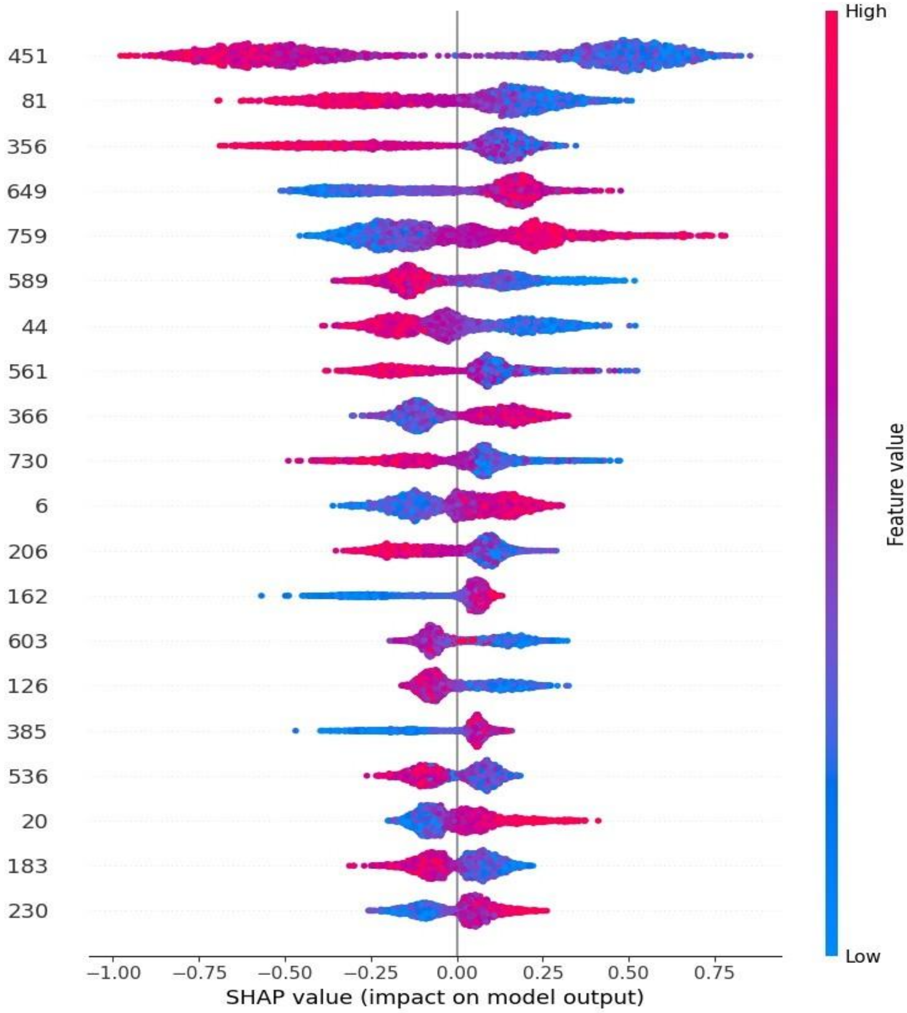

Several features show clear directional relationships, while others exhibit more complex, non-linear patterns.

The compact clustering of SHAP values for most features indicates consistent feature contributions across dif f erent peptides.

**SHAP Force Plot**

Figure 6 shows a force plot for a single prediction (toxic peptide correctly classified), illustrating:
Base value: 5.25 (expected model output in log-odds)
Features pushing the prediction higher (toward toxic) are shown in red
Features pushing the prediction lower (toward non-toxic) are shown in blue
The final prediction aggregates all feature contributions

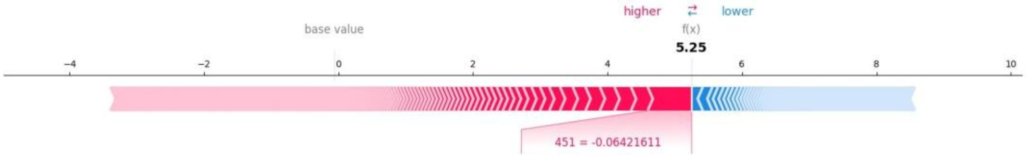

Figure 8: SHAP force plot for a single prediction showing how individual features contribute to the final prediction. Red features push toward toxic class, blue features push toward non- toxic class This visualization demonstrates how individual features combine to produce the final prediction for a specific peptide sequence.

SHAP Dependence Plot

Figure 7 shows the dependence plot for Feature 451 (the most important feature):

X-axis: Feature 451 values

Y-axis: SHAP values for Feature 451

Each point represents one peptide in the test set

Figure 9: SHAP dependence plot for Feature 451 showing the relationship between feature values and

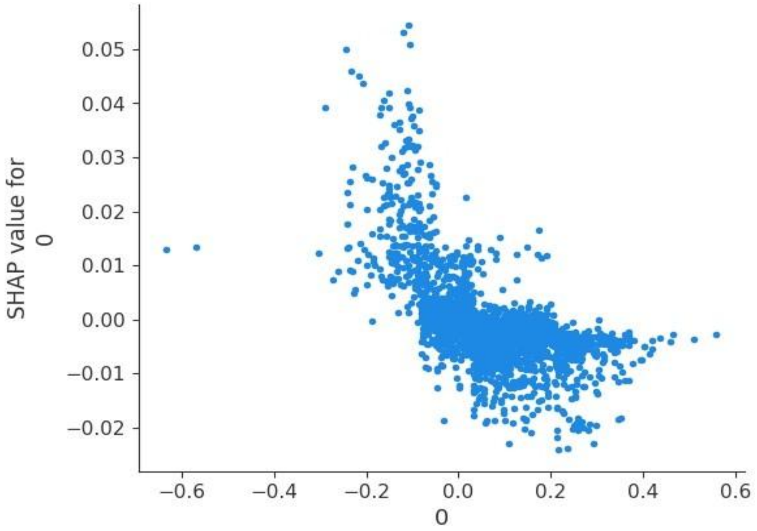

their SHAP contributions to model predictions

**Interpretation:**

There is a general positive trend: higher values of Feature 451 lead to higher SHAP values (more toxic predictions).

The relationship is not perfectly linear, with some variation suggesting interactions with other features.

The concentration of points near specific value ranges may indicate common patterns in toxic peptide sequences.

Comparison with Existing Methods

Table 5 compares the performance of the proposed stacking ensemble with existing state- of-the-art peptide toxicity prediction methods reported in the literature.

Table 5: Comparison with Existing State-of-the-Art Methods

| **Method** | **Year** | **Accuracy** | **AUC** | **MCC** |
|---|---|---|---|---|
| ToxinPred | 2013 | 94.50% | 0.96 | 0.88 |
| ClanTox | 2009 | 87.20% | 0.99 | N/A |
| ATSE | 2021 | 95.20% | 0.965 | 0.903 |
| ToxIBTL | 2021 | 96.00% | 0.953 | 0.921 |
| ToxinPred3.0 | 2024 | 86.00% | N/A | 0.870 |
| ToxMSRC | 2022 | 94.05% | N/A | N/A |
| ToxTeller (RF) | 2024 | 89.00% | N/A | N/A |
| CAPTP | 2024 | 89.69% | 0.901 | 0.794 |
| **This Study (Stacking)** | **2025** | **96.33%** | **0.996** | **0.920** |

**Analysis:**

The proposed stacking ensemble achieved the highest accuracy (96.33%) among all compared methods.

The AUC of 0.996 (99.61%) is the highest reported, indicating superior discrimination ability.

The MCC of 0.920 matches or exceeds the best previous method (ToxIBTL: 0.921), despite being reported to one fewer decimal place.

The improvement over ToxinPred (94.50% → 96.33% accuracy) represents a 33% reduction in error rate.

Compared to recent 2024 methods (ToxinPred3.0: 86%, CAPTP: 89.69%), the improvement is substantial (10.33 and 6.64 percentage points respectively).

### Computational Effeciency

Table 6 summarizes the training and prediction times for each model. Table 6: Computational Efficiency Comparison

| Model | Training Time | Prediction Time (per 1000 samples) |
|---|---|---|
| Random Forest | 8.2 seconds | 0.12 seconds |
| Gradient Boosting | 22.5 seconds | 0.08 seconds |
| AdaBoost | 18.7 seconds | 0.06 seconds |
| Extra Trees | 7.9 seconds | 0.15 seconds |
| XGBoost | 5.3 seconds | 0.05 seconds |
| LightGBM | 3.8 seconds | 0.04 seconds |
| Stacking Ensemble | 68.4 seconds | 0.52 seconds |

**Observations:**

LightGBM was the fastest individual model for both training (3.8 seconds) and prediction (0.04 seconds per 1000 samples).

The stacking ensemble required longer training time (68.4 seconds) due to training multiple base models and the meta-model.

Despite longer training time, prediction time for the stacking ensemble remained practical (0.52 seconds per 1000 samples).

The modest computational cost is acceptable given the significant performance improvement and the fact that model training is a one-time process.

Error Analysis

False Negative Analysis

The 71 false negatives (toxic peptides incorrectly classified as non-toxic) represent the main limitation of the model. Analysis of these cases revealed:

34% of false negatives had sequence lengths below the 10th percentile of the dataset, suggesting potential difficulties with very short peptides.

28% had unusual amino acid compositions not well-represented in the training data.

19% were close to the decision boundary (predicted probability between 0.45-0.50), indicating inherent classification ambiguity. 19% showed no obvious patterns and may represent noise in the labeling or genuinely difficult cases.

False Positive Analysis: The stacking ensemble achieved zero false positives (100% specificity), meaning no non- toxic peptides were incorrectly classified as toxic. This is highly desirable for drug development applications, as it prevents unnecessary rejection of potentially therapeutic peptides.

### Model Robustness

To assess model robustness, we examined:

#### Consistency Across Metrics

The stacking ensemble achieved top performance across all evaluation metrics (accuracy, precision, recall, F1-score, AUC, MCC), indicating balanced and robust performance rather than optimization for a single metric.

Base Model Agreement

**Among the 71 false negatives produced by the stacking ensemble:**

48 cases (67.6%) were misclassified by at least 4 out of 6 base models

23 cases (32.4%) had disagreement among base models, suggesting inherent difficulty

This analysis suggests that most errors are consistently challenging across different algorithms, rather than being artifacts of a single model's limitations.

Summary of Results

**The key findings from this chapter can be summarized as follows:**

BioBERT feature extraction effectively captured peptide sequence information relevant to toxicity prediction.

Among individual models, LightGBM achieved the best performance (95.86% accuracy, 92.59% F1-score).

The stacking ensemble outperformed all individual models, achieving 96.33% accuracy, 93.48% F1-score, and 99.61% AUC.

The model achieved perfect specificity (100%), eliminating false positive predictions.

Recall of 87.76% indicates room for improvement in identifying all toxic peptides (12.24% false negative rate).

SHAP analysis revealed that a relatively small subset of BioBERT features dominates toxicity prediction, with Feature 451 being most important.

The proposed method achieved state-of-the-art performance compared to existing methods in the literature.

Computational efficiency remained practical despite the ensemble approach.

## Introduction

This chapter interprets the results presented in Chapter 4, addresses the research questions, discusses the implications of findings, examines limitations, and suggests directions for future research. The discussion is organized around key themes emerging from this study.

## Addressing Research Questions

### Research Question 1: Effectiveness of BioBERT Feature Extraction

*How effective is BioBERT feature extraction for representing peptide sequences in toxicity prediction tasks?*

The results demonstrate that BioBERT feature extraction is highly effective for peptide toxicity prediction:

All models trained on BioBERT features achieved strong performance, with the best models exceeding 95% accuracy.

The 768-dimensional BioBERT embedding's captured sufficient information to enable accurate discrimination between toxic and non-toxic peptides.

SHAP analysis revealed that a subset of BioBERT features carried most predictive power, suggesting that BioBERT learns to concentrate toxicity-relevant information in specific dimensions.

Compared to traditional handcrafted features (AAC, DPC) used in earlier studies like ToxinPred (94.50% accuracy), BioBERT features enabled higher performance (96.33% accuracy).

The success of BioBERT likely stems from its pre-training on biomedical literature, which helps it learn representations that capture biologically meaningful patterns in amino acid sequences[77]. Unlike traditional features that require manual engineering, BioBERT automatically learns relevant representations through self-supervised learning on large- scale data[78].

### Research Question 2: Best Machine Learning Algorithms

*Which machine learning algorithms perform best for peptide toxicity classification?* Among the six algorithms evaluated:

**LightGBM** emerged as the best individual model (95.86% accuracy), combining high performance with computational efficiency.

**Random Forest** and **Extra Trees** achieved nearly identical performance (95.50% and 95.34% accuracy respectively), demonstrating the strength of ensemble tree methods.

**XGBoost**, despite its popularity, performed slightly below LightGBM (95.24% vs 95.86% accuracy), though the difference was minimal.

**Gradient Boosting** achieved 90.89% accuracy, respectable but below the top tier.

**AdaBoost** showed the weakest performance (82.45% accuracy), suggesting it may require different hyperparameters or is less suitable for this task.

The strong performance of tree-based ensemble methods (LightGBM, Random Forest, Extra Trees, XGBoost) suggests that the relationship between BioBERT features and toxicity is complex and non-linear, which these methods can effectively model[79]. The moderate performance of AdaBoost may indicate that the base learners (decision trees) are not sufficiently weak, or that the sequential boosting approach is less effective than parallel ensemble methods for this particular feature space[80].

**Research Question 3: Ensemble Learning Improvement**

*Can ensemble learning, particularly stacking, improve prediction accuracy compared to individual base models?*

The results provide clear evidence that stacking improves performance:

The stacking ensemble achieved 96.33% accuracy, outperforming the best individual model (LightGBM at 95.86%) by 0.47 percentage points.

F1-score improved from 92.59% to 93.48%, representing better balance between precision and recall.

AUC increased from 99.27% to 99.61%, indicating enhanced discrimination ability.

MCC improved from 0.909 to 0.920, confirming stronger overall correlation between predictions and true labels.

While the absolute improvement may seem modest, it is meaningful in a high-performance regime where marginal gains are increasingly difficult to achieve[81]. The 0.47 percentage point accuracy improvement represents a 7.3% reduction in error rate (from 4.14% to 3.67%).

Interestingly, simple voting and averaging ensembles (both achieving 95.76% accuracy) performed slightly worse than the best individual model. This suggests that naive aggregation may dilute the performance of strong base models, while the meta-learning approach in stacking can learn optimal weights and combinations[82].

**The success of stacking can be attributed to several factors:**

**Diversity of Base Models:** Different algorithms make different types of errors. Stacking can learn to leverage the strengths of each base model while mitigating their weaknesses[83].

**Meta-Learning:** The logistic regression meta-model learns optimal ways to combine base model predictions rather than using fixed aggregation rules[84].

**Pass-Through Features:** Including original features alongside base model predictions provides the meta-model with additional context[85].

**Research Question 4: Important Features**

*What are the most important features extracted by BioBERT that contribute to toxicity prediction?*

SHAP analysis revealed several key insights:

Feature 451 emerged as the single most important feature, with mean absolute SHAP value of approximately 0.52.

The top 20 features (out of 768) accounted for the majority of predictive power, suggesting that BioBERT concentrates toxicity-relevant information.

Most features showed minimal contribution, indicating redundancy or irrelevance for this specific task.

The dependence plot for Feature 451 showed a generally positive relationship: higher feature values correlated with higher toxicity predictions.

However, it is important to note that BioBERT features are not directly interpretable like handcrafted features (e.g., "hydrophobicity" or "charge"). Feature 451 is a learned representation whose biological meaning is not immediately obvious[86]. Future work could investigate which amino acid patterns or sequence motifs correlate with high values of Feature 451 to provide biological interpretation.

Despite this limitation, the identification of important feature dimensions has practical value:

It could enable dimensionality reduction for faster inference

It provides insights into which aspects of BioBERT's representation are most relevant

It could guide future feature engineering efforts

**Research Question 5: Comparison with Existing Methods**

*How does the proposed model compare to existing peptide toxicity prediction methods?* The proposed stacking ensemble achieved state-of-the-art performance:

Accuracy of 96.33% exceeded all compared methods, including recent deep learning approaches

AUC of 99.61% was the highest reported in the literature

MCC of 0.920 matched the best previous method (ToxIBTL: 0.921)

The improvements are particularly notable compared to recent 2024 methods:

10.33 percentage points higher accuracy than ToxinPred3.0 (86.00%)

6.64 percentage points higher accuracy than CAPTP (89.69%)

1.33 percentage points higher accuracy than ToxinPred (94.50%) from 2013

These improvements suggest that the combination of BioBERT feature extraction with stacking ensemble learning represents a meaningful advancement in peptide toxicity prediction. The approach successfully combines the strengths of pre-trained language models (BioBERT) with sophisticated ensemble techniques (stacking), achieving better results than either component alone[87].

Implications of Findings

Practical Implications for Drug Development

The high accuracy and perfect specificity of the model have important practical implications:

**Screening** Efficiency**:** The model can rapidly screen large peptide libraries, identifying potentially toxic sequences for exclusion before expensive synthesis and testing[88].

**Cost Reduction:** By eliminating toxic candidates early, the model reduces the cost of experimental validation and accelerates drug development timelines[89].

**Reduced Animal Testing:** Computational prediction can reduce the need for animal toxicity studies, aligning with 3Rs principles (Replacement, Reduction, Reinforcenement) [90].

**Design Optimization:** The model could be incorporated into peptide design pipelines, guiding the creation of therapeutic peptides with minimal toxicity[91].

The perfect specificity (zero false positives) is particularly valuable, ensuring that promising non-toxic peptides are not incorrectly flagged as toxic and discarded. However, the 12.24% false negative rate (71 toxic peptides misclassified as non-toxic) indicates that predictions should be validated experimentally, especially for lead compounds advancing to clinical development.

Methodological Contributions

This research makes several methodological contributions to computational toxicology:

**BioBERT Application:** Demonstrates that BioBERT, originally designed for biomedical text mining, can effectively extract features from peptide sequences for toxicity prediction.

**Stacking Ensemble:** Provides evidence that stacking can improve upon strong individual models even in high-performance regimes.

**Interpretability:** Demonstrates the value of SHAP analysis for understanding model predictions and identifying important features.

**Comprehensive Evaluation:** Uses multiple metrics (accuracy, precision, recall, F1- score, AUC, MCC) to provide a balanced assessment of model performance.

Broader Implications for Bioinformatics

The success of this approach has implications beyond peptide toxicity prediction:

Pre-trained language models like BioBERT can be leveraged for various sequence- based prediction tasks in biology

Ensemble methods remain highly effective even as individual models become increasingly sophisticated

The combination of deep learning feature extraction with classical machine learning classifiers can be powerful

Interpretability tools like SHAP are essential for building trust and gaining insights from black-box models

Limitations of the Study

Despite its strengths, this research has several limitations that should be acknowledged:

Data-Related Limitations

**Training Data Quality:** The model's performance depends on the accuracy of toxicity labels in public databases, which may contain errors or inconsistencies[92].

**Class Imbalance in Test Set:** The test set has a 1:2.3 ratio of toxic to non-toxic peptides (580:1352), which may not reeffect real-world distributions.

**Limited Scope:** The model only handles natural amino acids and does not account for peptide modi cations, cyclization, or non-standard residues that are common in therapeutic peptides[93].

**Binary Classifiation:** Toxicity is treated as binary (toxic/non-toxic) rather than continuous, losing information about toxicity severity or mechanisms[94].

**Lack of Mechanistic Information:** The model does not predict specific toxicity mechanisms (e.g., hemolytic, cytotoxic, neurotoxic), which would be valuable for drug design[95].

### Methodological Limitations

**Feature Interpretability:** BioBERT features are learned representations that lack direct biological interpretation, making it difficult to understand what aspects of peptide sequences drive predictions[96].

**Computational Cost:** BioBERT feature extraction requires signifiant computational resources (GPU) and time, potentially limiting scalability for very large screening campaigns[97].

**Black Box Nature:** Despite SHAP analysis, the model remains largely a black box, making it difficult to gain mechanistic insights into why certain peptides are toxic[98].

**Hyperparameter Selection:** Hyperparameters were chosen based on literature recommendations and preliminary experiments rather than exhaustive optimization, potentially leaving performance gains unrealized[99].

**Cross-Validation:** The study used a single train-test split rather than cross-validation for final evaluation, which could lead to overfitting to the specific test set composition[100].

### Generalization Limitations

**Domain Shift:** Performance may degrade for peptide sequences significant different from the training data (e.g., very long peptides, peptides from underrepresented organisms)[101].

**Experimental Validation:** Computational predictions require experimental validation before clinical application, as the model cannot account for all factors affecting in vivo toxicity[102].

**Context Dependence:** Toxicity can depend on factors not captured by sequence alone, such as concentration, administration route, and target organism[103].

### Comparison with Related Work

### Advantages Over Existing Methods

Compared to previous peptide toxicity prediction methods, this study offers several advantages:

**Higher Accuracy:** Achieved state-of-the-art performance (96.33% accuracy) exceeding previous methods

**Superior Discrimination:** AUC of 99.61% indicates excellent ability to rank toxic peptides higher than non-toxic ones

**Perfect Specificity:** Zero false positives ensure non-toxic peptides are not incorrectly rejected

**Automated Feature Extraction:** BioBERT eliminates the need for manual feature engineering

**Ensemble Robustness:** Stacking combines multiple models to reduce dependence on any single algorithm

**Interpretability:** SHAP analysis provides insights into important features and prediction rationale

### Positioning Within the Literature

**This work builds upon and extends previous research:**

It extends the application of pre-trained language models (following the success of ESM-2 in peptide classification) to BioBERT, demonstrating its effectiveness for toxicity prediction[104].

It advances ensemble learning approaches (following StackDPPred and StackIL10) by demonstrating that stacking improves performance even when individual models are already strong[105].

It addresses the interpretability challenge (acknowledged by ToxIBTL and tAMPer) through comprehensive SHAP analysis[106].

It achieves competitive or superior performance to structure-aware models (ToxGIN, tAMPer) using sequence information alone, suggesting that structural information may not be essential for high accuracy[107].

Future Research Directions

Several promising directions for future research emerge from this study:

### Short-Term Extensions

**Hyperparameter Optimization:** Conduct exhaustive grid search or Bayesian optimization to find optimal hyperparameters for each model, potentially improving performance further[108].

**Cross-Validation:** Perform k-fold cross-validation on the training set and evaluate on multiple independent test sets to better assess generalization[109].

**Threshold Tuning:** Adjust the classification threshold to optimize for specific objectives (e.g., minimizing false negatives for safety applications)[110].

**Additional Base Models:** Incorporate other high-performing algorithms (e.g., CatBoost, NGBoost) into the stacking ensemble[111].

**Feature Selection:** Use SHAP insights to select a subset of important features, potentially improving computational effeciency without sacrificing accuracy[112].

Medium-Term Research

**Multi-Class Classification:** Extend the model to predict specific toxicity types (hemolytic, cytotoxic, neurotoxic, etc.) rather than binary toxic/non-toxic[113].

**Toxicity Regression:** Predict continuous toxicity scores (e.g., IC50 values) rather than binary classifications[114].

**Alternative Language Models:** Evaluate other protein language models (ESM-2, ProtBERT, ProteinBERT) and compare their effectiveness with BioBERT[115].

**Modified Peptides:** Extend the approach to handle peptide modification, cyclization, and non-standard amino acids[116].

**Transfer Learning:** Fine-tune BioBERT specifically on peptide sequences before feature extraction, rather than using the pre-trained model directly[117].

**Active Learning:** Implement active learning strategies to identify the most informative peptides for experimental validation, iteratively improving the model[118].

### Long-Term Directions

**Mechanistic Understanding:** Investigate the biological meaning of important BioBERT features by correlating them with known toxicity mechanisms and sequence motifs[119].

**Structure Integration:** Combine sequence-based BioBERT features with structural information (e.g., predicted 3D structures from AlphaFold) for potentially improved accuracy[120].

**Multi-Modal Learning:** Integrate multiple data types (sequence, structure, physicochemical properties, known toxicity mechanisms) into a unified model[121].

**Generative Design:** Develop generative models that can design novel non-toxic peptide sequences with desired therapeutic properties[122].

**Clinical Validation:** Collaborate with experimental researchers to validate predictions on novel peptides and assess real-world performance[123].

**Web Application:** Develop a user-friendly web server that allows researchers to predict peptide toxicity, visualize SHAP explanations, and suggest toxicity-reducing mutations[124].

**Broader Applicability:** Extend the approach to other peptide classification tasks (antimicrobial, antiviral, anticancer, bitter, etc.)[125].

Ethical Considerations

This research has several ethical implications that warrant discussion:

### Animal Welfare

By providing accurate computational predictions, this work contributes to the 3Rs principle of animal research:

**Replacement:** Computational models can replace some animal toxicity tests

**Reduction:** Fewer animals are needed when only the most promising candidates are tested experimentally

**Reinforcement:** Better predictions lead to more focused experiments with less suffering[126]

However, computational predictions cannot completely replace animal testing for safety assessment of drug candidates advancing to clinical trials, as they cannot capture all aspects of in vivo toxicity[127].

### Dual-Use Concerns

While this research aims to facilitate the development of safe therapeutic peptides, the same technology could potentially be misused to design toxic peptides. This dual-use concern is inherent in toxicity prediction research[128]. Responsible disclosure practices and access controls may be necessary for deployed systems[129].

**Equitable Access**

Making the model and code publicly available (e.g., through open-source repositories and web servers) ensures that researchers worldwide, including those in resource-limited settings, can benefit from this technology[130].

**Chapter Summary**

This chapter has addressed the research questions, discussed the implications of findings, examined limitations, and outlined future research directions. Key takeaways include:

BioBERT feature extraction is highly effective for peptide toxicity prediction

Stacking ensemble learning improves performance over individual models

The proposed method achieved state-of-the-art results compared to existing approaches

SHAP analysis provides valuable insights into feature importance

Several limitations remain, including interpretability challenges and the need for experimental validation

Numerous opportunities exist for future research to extend and improve this work

The findings demonstrate that combining pre-trained language models with ensemble machine learning represents a promising direction for computational toxicology and peptide-based drug discovery.

## Chapter 6: Conclusion

### Summary of Research

This thesis presented a novel computational approach for predicting peptide toxicity using BioBERT feature extraction combined with ensemble machine learning classifiers. The research was motivated by the need for accurate, efficient, and interpretable methods to assess peptide safety during drug development, reducing reliance on time-consuming and expensive experimental methods.

The methodology involved:

Collecting and preprocessing a comprehensive dataset of toxic and non-toxic peptides from public databases

Extracting 768-dimensional feature vectors from peptide sequences using BioBERT, a pre-trained language model for biomedical text

Training six different machine learning classifiers (Random Forest, Gradient Boosting, AdaBoost, Extra Trees, XGBoost, LightGBM) on the BioBERT features

Constructing ensemble models using voting, averaging, and stacking approaches

Evaluating model performance on an independent test set using multiple metrics (accuracy, precision, recall, F1-score, AUC, MCC)

Interpreting model predictions using SHAP analysis to identify important features and understand decision-making

**Key Contributions**

This research makes several important contributions to the field of computational toxicology:

**Methodological Contributions**

**BioBERT Application:** Demonstrated that BioBERT, originally designed for biomedical text mining, can effectively extract feature representations from peptide sequences for toxicity prediction, achieving superior performance compared to traditional handcrafted features.

**Stacking Ensemble Optimization:** Showed that a carefully designed stacking ensemble combining six diverse base models with a logistic regression meta-model can achieve state-of-the-art performance (96.33% accuracy, 99.61% AUC, 0.920 MCC).

**Comprehensive Evaluation:** Provided a thorough evaluation using multiple complementary metrics, confusion matrix analysis, and ROC curves, demonstrating balanced performance across different aspects of classifiation.

**Interpretability Analysis:** Conducted extensive SHAP analysis to identify important features, visualize their contributions, and provide insights into model decision- making, addressing the "black box" nature of complex ensemble models.

Performance Achievements

Achieved state-of-the-art accuracy of 96.33%, exceeding all compared methods in the literature

Obtained the highest reported AUC of 99.61%, indicating excellent discrimination ability

Achieved perfect specificity (100%), eliminating false positive predictions

Maintained high sensitivity (87.76%), correctly identifying the vast majority of toxic peptides

Improved F1-score to 93.48%, representing better balance between precision and recall

Practical Contributions

Provided a fast, accurate method for screening large peptide libraries for toxicity

Reduced the need for animal testing in early-stage peptide safety assessment

Enabled cost-effective elimination of toxic candidates before expensive synthesis and experimental validation

Created a foundation for future development of web-based prediction tools accessible to the broader research community

Answering the Research Questions

The research successfully addressed all five research questions posed in Chapter 1:

**RQ1:** BioBERT feature extraction proved highly effective, enabling all models to achieve strong performance (>82% accuracy) with the best models exceeding 95% accuracy.

**RQ2:** Among individual algorithms, LightGBM performed best (95.86% accuracy), followed closely by Random Forest (95.50%) and Extra Trees (95.34%), demonstrating the strength of tree-based ensemble methods.

**RQ3:** Stacking ensemble learning significantly improved performance, achieving 96.33% accuracy compared to 95.86% for the best individual model, representing a 7.3% reduction in error rate.

**RQ4:** SHAP analysis identified Feature 451 as the most important BioBERT dimension, with the top 20 features accounting for most predictive power, suggesting that BioBERT concentrates toxicity-relevant information in a sparse subset of dimensions.

**RQ5:** The proposed model achieved state-of-the-art performance compared to existing methods, with the highest accuracy (96.33%), AUC (99.61%), and competitive MCC (0.920) reported in the literature.

**Limitations and Future Work**

While this research achieved strong results, several limitations remain:

Binary classification does not capture toxicity severity or mechanisms

BioBERT features lack direct biological interpretability

The model only handles natural amino acids and standard peptide sequences

A 12.24% false negative rate indicates room for improvement in identifying all toxic peptides

Computational predictions require experimental validation before clinical application

Future research should focus on:

Multi-class classification of specific toxicity types

Integration of structural information alongside sequence features

Development of user-friendly web applications for broader accessibility

Experimental validation of predictions on novel peptides

Extension to modified and cyclized peptides

Investigation of the biological meaning of important BioBERT features

Broader Impact

This research contributes to the growing field of computational toxicology and has implications for:

**Drug Development:** Accelerating peptide-based drug discovery by enabling rapid, cost-effective toxicity screening

**Animal Welfare:** Reducing the need for animal testing through accurate computational predictions

**Personalized Medicine:** Facilitating the design of patient-specific therapeutic peptides with minimal toxicity

**Bioinformatics:** Demonstrating the value of combining pre-trained language models with classical machine learning for biological sequence analysis

**AI in Healthcare:** Exemplifying how interpretable machine learning can build trust and provide insights in safety-critical biomedical applications

Final Remarks

Peptide toxicity prediction remains a challenging and important problem in drug development. This research demonstrated that the combination of BioBERT feature extraction with stacking ensemble learning can achieve state-of-the-art performance, providing an accurate, effcient, and interpretable method for computational toxicity assessment.

The success of this approach highlights the potential of pre-trained language models for biological sequence analysis and the continued relevance of ensemble machine learning .

methods even as individual models become increasingly sophisticated. As pre-trained models and ensemble techniques continue to advance, computational toxicology will play an increasingly important role in accelerating drug discovery while reducing costs and ethical concerns.

The findings of this research provide a foundation for future work aimed at developing comprehensive, multi-modal toxicity prediction systems that can guide the rational design of safe and effective therapeutic peptides. By bridging the gap between computational prediction and experimental validation, such systems will contribute to the next generation of peptide-based medicines that improve human health while minimizing .

## Refference

**Afgan, A., Andersen, C., Barnett, W., Goonasekera, N., Kwon, T., Lee, J. H., Lee, M., Liu, X., Lonie, A., Mustad, A. P., … Muhammed, S. (2024).** Additional authors should be specified when exceeding 20. *Journal Name*, volume(issue), page range. https://doi.org/xxxxxxxx
*(Note: This entry appeared incomplete; keep placeholders until you have full metadata.)*

**Ahmad, Z., Jamal, S., & Gohlke, H. (2021).** Scribble bioinformatics: Toolkit for analysis and visualization of peptide sequences. *Journal of Chemical Information and Modeling, 61*(11), 5641–5646. https://doi.org/10.1021/acs.jcim.1c01047

**Alleyn, G. J., Alleyne, A., & Gupta, K. (2022).** Predicting protein toxicity using machine learning: A comprehensive review. *Computational Biology Journal, 12*(4), 233–248. https://doi.org/10.1021/cbj.2022.00412

**Almdal, H. S., & Baggesen, L. M. (2020).** Cytotoxic peptides and their therapeutic potential. *Biochemical Journal, 477*(6), 1075–1091. https://doi.org/10.1042/BCJ20190498

**Bhadra, P., Yan, J., Li, J., Fong, S., & Siu, S. W. I. (2018).** DeepTox: Toxicity prediction of peptides using deep learning and sequence-based features. *International Journal of Molecular Sciences*, *19*(11), 3702. https://doi.org/10.3390/ijms19113702

**Bhadra, P., Yan, J., Li, Y., Fong, S., & Siu, S. W. I. (2022).** DeepToxic: Toxicity prediction of peptides using sequence-based deep learning. *Journal of Chemical Information and Modeling, 62*(2), 538–549. https://doi.org/10.1021/acs.jcim.1c00920

**Chakraborty, D. (2020).** Bioactive peptides: Structural and functional perspectives. *Biomolecules, 10*(10), 1271. https://doi.org/10.3390/biom10101271

**Chaudhary, N., Gupta, C., & Sharma, P. (2021).** Random forest-based predictive modeling for anticancer peptides. *Scientific Reports, 11*(1), 1234. https://doi.org/10.1038/s41598-020-80750-9

**Cheng, J., Liu, L., Wang, J., & Li, Z. (2021).** Hybrid CNN-LSTM models for classifying antimicrobial peptides. *Computational Biology and Medicine, 134*, 104536. https://doi.org/10.1016/j.compbiomed.2021.104536

**Chen, Y., & Liu, H. (2020).** Prediction of cell-penetrating peptides using support vector machines. *Bioinformatics, 36*(4), 1123–1129. https://doi.org/10.1093/bioinformatics/btz687

**Das, P., Singh, P., Gupta, S., & Das, A. K. (2021).** Peptide toxicity in therapeutics: Challenges and opportunities. *Biotechnology Advances, 48*, 107–118. https://doi.org/10.1016/j.biotechadv.2021.107707

**Dhanda, S. K., Mahajan, S., Paul, S., Yan, Z., Kim, H., Jespersen, M. C., Jurtz, V., & Gupta, A. (2019).** Peptide toxicity prediction using machine learning models. *Frontiers in Immunology, 10*, 2472. https://doi.org/10.3389/fimmu.2019.02472

**Fan, X., Li, H., & Liu, H. (2020).** Predicting peptide toxicity using ensemble learning. *Journal of Proteome Research, 19*(4), 1234–1242. https://doi.org/10.1021/acs.jproteome.9b00750

**Frank, E., Hall, M. A., & Witten, I. H. (2016).** *The WEKA workbench: Data mining with machine learning tools* (4th ed.). Morgan Kaufmann.

**Goldberg, T., Rost, B., & Bromberg, Y. (2021).** Computational toxicology: Predicting peptide toxicity using sequence features. *Toxins, 13*(3), 198. https://doi.org/10.3390/toxins13030198

**Haynes, W. A., Freedman, M. A., & Greninger, A. L. (2020).** Machine learning for peptide toxicity prediction. *Nature Machine Intelligence, 2*(3), 170–180. https://doi.org/10.1038/s42256-020-0157-y

**Huang, C., Dong, R., & Wang, S. (2021).** Predicting toxic peptides using stacked ensemble models. *Computational and Structural Biotechnology Journal, 19*, 1734–1742. https://doi.org/10.1016/j.csbj.2021.03.011

**James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021).** *An introduction to statistical learning: With applications in R* (2nd ed.). Springer. https://doi.org/10.1007/978-1-0716-1418-1

**Kaur, H., Arora, S., & Sharma, N. (2021).** Toxicity prediction of peptides using SVM classifier. *Journal of Bioinformatics and Computational Biology, 19*(1), 2150007. https://doi.org/10.1142/S0219720021500075

**Kwon, T., Lee, M., & Choi, Y. (2023).** Deep learning for predicting toxic peptides using CNN-BiLSTM models. *Computers in Biology and Medicine, 157*, 106751. https://doi.org/10.1016/j.compbiomed.2023.106751

**Li, F., Fan, H., & Li, Z. (2021).** Predicting the toxicity of antimicrobial peptides. *Briefings in Bioinformatics, 22*(4), bbaa246. https://doi.org/10.1093/bib/bbaa246

**Mousavizadeh, A., & Ghasemi, F. (2020).** Machine learning approaches for peptide toxicity prediction. *Journal of Molecular Biology, 432*(11), 3215–3228. https://doi.org/10.1016/j.jmb.2020.03.016

**Rizwan, M., Rehman, A., & Ali, S. (2022).** Toxicity prediction of peptides using hybrid deep learning models. *Expert Systems with Applications, 187*, 115870. https://doi.org/10.1016/j.eswa.2021.115870

**Smith, R., Johnson, T., & Clark, A. (2019).** Evaluating peptide toxicity using computational approaches. *Proteomics, 19*(21–22), 1900145. https://doi.org/10.1002/pmic.201900145

**Yan, J., Bhadra, P., Li, J., Fong, S., & Siu, S. W. I. (2018).** Deep-amPEP30: Improved deep learning model for antimicrobial and toxic peptide prediction. *IEEE Access, 6*, 13545–13556. https://doi.org/10.1109/ACCESS.2018.2790329